

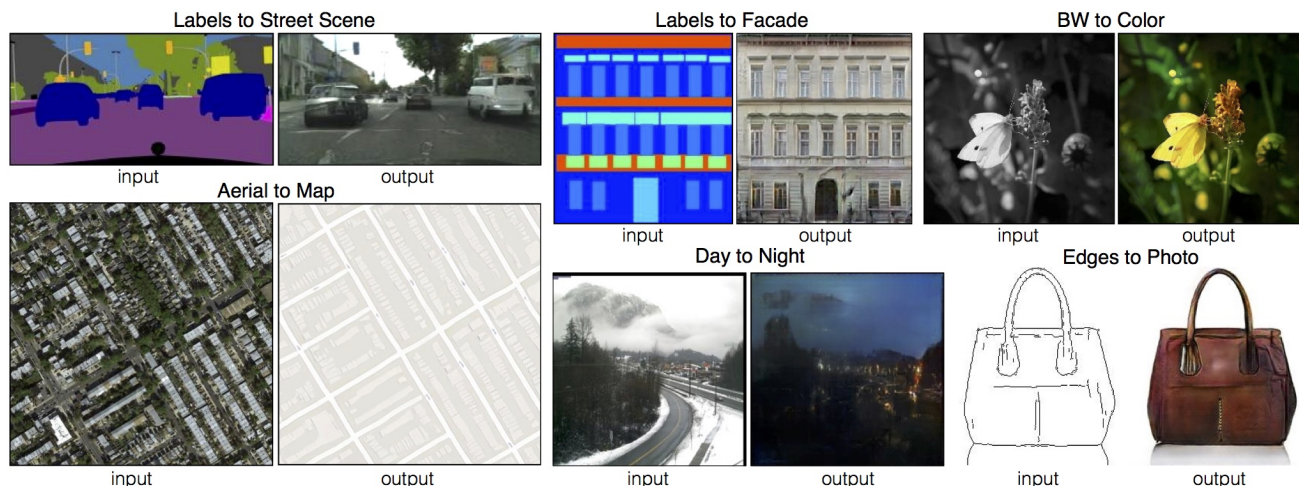
Image-to-Image Translation with Conditional Adversarial Nets

Phillip Isola

Jun-Yan Zhu

Tinghui Zhou

Alexei A. Efros

 University of California, Berkeley
 In CVPR 2017
[\[Paper\]](#)[\[GitHub\]](#)

Example results on several image-to-image translation problems. In each case we use the same architecture and objective, simply training on different data.

Abstract

We investigate conditional adversarial networks as a general-purpose solution to image-to-image translation problems. These networks not only learn the mapping from input image to output image, but also learn a loss function to train this mapping. This makes it possible to apply the same generic approach to problems that traditionally would require very different loss formulations. We demonstrate that this approach is effective at synthesizing photos from label maps, reconstructing objects from edge maps, and colorizing images, among other tasks. As a community, we no longer hand-engineer our mapping functions, and this work suggests we can achieve reasonable results without hand-engineering our loss functions either.

Try our code

 Recommended version: [\[PyTorch\]](#)

 Original code: [\[Torch\]](#)

Ports of our code:

[\[Tensorflow\]](#) (tutorial on Tensorflow Core)

[\[Tensorflow\]](#) (implementation by [Christopher Hesse](#))


[\[Tensorflow\]](#) (implementation by [Yen-Chen Lin](#))

[\[Chainer\]](#) (implementation by [pfnet-research](#))

[\[Keras\]](#) (implementation by [Thibault de Boissiere](#))

[\[Wolfram Cloud\]](#) (implementation by Wolfram team)

Paper




Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros
Image-to-Image Translation with Conditional Adversarial Networks
 CVPR, 2017 ([Paper](#))

[\[bibtex\]](#)


Interactive Demo

(made by [Christopher Hesse](#))



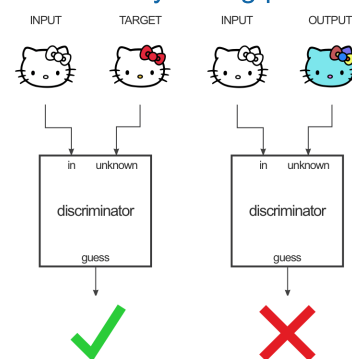
Expository articles and videos

Two-minute Papers



Karoly Zsolnai-Feher made the above as part of his very cool ["Two-minute papers" series](#).

Affinelayer blog post



Great explanation by [Christopher Hesse](#), also documenting his [tensorflow port](#) of our code.

Experiments

Here we show comprehensive results from each experiment in our paper. Please see the paper for details on these experiments.

Effect of the objective

[Cityscapes](#)
[Facades](#)

Effect of the generator architecture

[Cityscapes](#)

Effect of the discriminator patch scale

[Cityscapes](#)
[Facades](#)


Additional results

[Map to aerial](#)
[Aerial to map](#)
[Semantic segmentation](#)
[Day to night](#)
[Edges to handbags](#)
[Edges to shoes](#)
[Sketches to handbags](#)
[Sketches to shoes](#)

Community contributions: [#pix2pix](#)


People have used our code for many creative applications, often posted on twitter with the hashtag [#pix2pix](#). Check them out [here](#)! Below we highlight just a few of the many:

Human Allocation of Space




[Scott Eaton](#) uses a customized version of pix2pix as a tool in his artwork, for example training a net to [translate from](#)

Learning to See: Gloomy Sunday



[Memo Akten](#) used pix2pix to create the very compelling music video linked above, in which common household items, like a powercord, are moved around in a pantomime of crashing waves and blooming flowers. Then a pix2pix-based model translates the

Suggestive Drawing



[Nono Martinez Alonso](#) used pix2pix for his masters thesis on human-AI collaboration for design. Different pix2pix models are used as aids for a human designer, where they "suggest" extensions, refinements, and completions of the designer's sketches.

sketches and brush strokes to 3D renderings. The sculpture above is an actual bronze cast derived from one of Scott's translated designs. You can find more of his AI-empowered artwork [here](#).

#edges2cats



Christopher Hesse trained our model on converting edge maps to photos of cats, and included this in his [interactive demo](#). Apparently, this is what the Internet wanted most, and #edges2cats briefly [went viral](#). The above cats were designed by Vitaly Vidmironov ([@vvid](#)).

Interactive Anime



Bertrand Gondouin trained our method to translate sketches→Pokemon, resulting in an interactive drawing tool.

pantomime into renderings of the imagined objects.

Alternative Face



Mario Klingemann used our code to translate the appearance of French singer Francoise Hardy onto Kellyanne Conway's infamous "alternative facts" interview. Interesting articles about it can be read [here](#) and [here](#).

Person-to-Person




Brannon Dorsey recorded himself mimicking frames from a video of Ray Kurzweil giving a talk. He then used this data to train a Dorsey→Kurzweil translator, allowing him to become a kind of puppeteer in control of Kurzweil's appearance.

Background masking



Kaihu Chen performed a number of [interesting experiments](#) using our method, including getting it to mask out the background of a portrait as shown above.

Color palette completion



Colormind adapted our code to predict a complete 5-color palette given a subset of the palette as input. This application stretches the definition of what counts as "image-to-image translation" in an exciting way: if you can visualize your input/output data as images, then image-to-image methods are applicable! (not that this is necessarily the best choice of representation, just one to think about.)

Recent Related Work

Generative adversarial networks have been vigorously explored in the last two years, and many conditional variants have been proposed. Please see the discussion of related work in [our paper](#). Below we point out three papers that especially influenced this work: the original GAN paper from Goodfellow et al., the [DCGAN framework](#), from which our code is derived, and the [iGAN](#) paper, from our lab, that first explored the idea of using GANs for mapping use strokes to images.

Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. **Generative Adversarial Networks**. NIPS, 2014. [[PDF](#)]

Alec Radford, Luke Metz, Soumith Chintala. **Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks**. ICLR, 2016. [[PDF](#)][[Code](#)]

Jun-Yan Zhu, Philipp Krahenbuhl, Eli Shechtman, Alexei A. Efros. **Generative Visual Manipulation on the Natural Image Manifold**. ECCV, 2016. [[PDF](#)][[Webpage](#)][[Code](#)]

Also, please check out our follow-up work on image-to-image translation *without* paired training examples:

Jun-Yan Zhu*, Taesung Park*, Phillip Isola, Alexei A. Efros. **Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks**. ICCV, 2017. [[PDF](#)][[Webpage](#)][[Code](#)]

Acknowledgements

We thank Richard Zhang, Deepak Pathak, and Shubham Tulsiani for helpful discussions. Thanks to Saining Xie for help with the HED edge detector. Thanks to the online community for exploring many applications of our work and pointing out typos and errors in the paper and code. This work was supported in part by NSF SMA-1514512, NGA NURI, IARPA via Air Force Research Laboratory, Intel Corp, Berkeley Deep Drive, and hardware donation by Nvidia. Disclaimer: The views and conclusions contained herein are those of the authors and should not be interpreted as necessarily representing the official policies or endorsements, either expressed or implied, of IARPA, AFRL or the U.S. Government.